

Suggested environment: Ubuntu 20 LTS, ansible 2.9.16,

```
root@ansible: ~  
root@ansible:~# lsb_release -a  
No LSB modules are available.  
Distributor ID: Ubuntu  
Description:    Ubuntu 20.04.6 LTS  
Release:        20.04  
Codename:       focal  
root@ansible:~#
```

ansible 2.9.16

```
root@ansible: ~  
root@ansible:~# ansible --version  
ansible 2.9.16  
  config file = None  
  configured module search path = ['/root/.ansible/plugins/modules', '/usr/share/ansible/plugins/modules']  
  ansible python module location = /usr/local/lib/python3.8/dist-packages/ansible  
  executable location = /usr/local/bin/ansible  
  python version = 3.8.10 (default, Sep 11 2024, 16:02:53) [GCC 9.4.0]  
root@ansible:~#
```

1) Which ansible command can display all ansible_ configuration for a host.

```
#ansible all -m setup
```

1) Please configure a cron job that runs logrotate on all machines every 10 minutes between 2h - 4h.

Ansible Playbook:

```
- hosts:

    - app-vm1.fra1.internal

    - db-vm1.fra1.db

    - web-vm1.fra1.web

tasks:

    - name: Install ntpd

      apt:

        name: ntp

        state: present

    - name: Deploy custom ntpd.conf

      copy:

        src: /home/dnet/ntpd.conf # Ensure this points to the actual
file
        dest: /etc/ntp.conf

        owner: root

        group: root

        mode: '0644'
```

```

- hosts:
  - all
tasks:
  - name: Configure logrotate cron job
    cron:
      name: "Run logrotate every 10 minutes between 2 AM and 4 AM"
      minute: "*/10"
      hour: "2-3" # 2 AM to 3:59 AM
      job: "/usr/sbin/logrotate /etc/logrotate.conf"
      state: present

```

```

root@ansible:~/ansible-playbooks# ansible-playbook configure_cron.yml -l servers

PLAY [all] *****

TASK [Gathering Facts] *****
ok: [web-vm1.fral.web]
ok: [app-vm1.fral.internal]
ok: [db-vm1.fral.db]

TASK [Configure logrotate cron job] *****
changed: [web-vm1.fral.web]
changed: [db-vm1.fral.db]
changed: [app-vm1.fral.internal]

PLAY RECAP *****
app-vm1.fral.internal      : ok=2    changed=1    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
db-vm1.fral.db             : ok=2    changed=1    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
web-vm1.fral.web           : ok=2    changed=1    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0

root@ansible:~/ansible-playbooks#

```

Its set to every server

```

root@app-vm1: ~
root@app-vm1:~# crontab -l
#Ansible: Run logrotate every 10 minutes between 2 AM and 4 AM
*/10 2-3 * * * /usr/sbin/logrotate /etc/logrotate.conf
root@app-vm1:~#

```

```
#ansible-playbook playbook.yml -l servers
```

```
root@ansible:~# ansible-playbook playbook.yml -l servers

PLAY [app-vm1.fra1.internal,db-vm1.fra1.db,web-vm1.fra1.web] *****

TASK [Gathering Facts] *****
ok: [web-vm1.fra1.web]
ok: [app-vm1.fra1.internal]
ok: [db-vm1.fra1.db]

TASK [Install ntpd] *****
ok: [web-vm1.fra1.web]
ok: [app-vm1.fra1.internal]
ok: [db-vm1.fra1.db]

TASK [Deploy custom ntpd.conf] *****
changed: [web-vm1.fra1.web]
changed: [app-vm1.fra1.internal]
changed: [db-vm1.fra1.db]

PLAY RECAP *****
app-vm1.fra1.internal      : ok=3    changed=1    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
db-vm1.fra1.db            : ok=3    changed=1    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
web-vm1.fra1.web          : ok=3    changed=1    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0

root@ansible:~# █
```

We also need to deploy monitoring template onto our nagios server
“monitoring.fra1.internal”, each of the above machines should use the following nagios templates:

```
# Host Definitions
define host {
    host_name                app-vm1.fra1.internal
    address                  192.168.0.2
    check_command             check_ping
    active_checks_enabled    1
    passive_checks_enabled   1
    max_check_attempts        3
}

define host {
    host_name                db-vm1.fra1.db
    address                  192.168.0.3
    check_command             check_ping
```

```

    active_checks_enabled      1
    passive_checks_enabled     1
    max_check_attempts         3
}

define host {
    host_name                  web-vm1.fra1.web
    address                    192.168.0.4
    check_command               check_ping
    active_checks_enabled      1
    passive_checks_enabled     1
    max_check_attempts         3
}

# Service Definitions
define service {
    service_description        ntp_process
    host_name                   app-vm1.fra1.internal
    check_command               check_ntp
    check_interval              10
    max_check_attempts          3
}

define service {
    service_description        ntp_process
    host_name                   db-vm1.fra1.db
    check_command               check_ntp
    check_interval              10
    max_check_attempts          3
}

define service {
    service_description        ntp_process
    host_name                   web-vm1.fra1.web
    check_command               check_ntp
    check_interval              10
    max_check_attempts          3
}

```

← → ↻ Not secure monitoring.fra1.internal/

Nagios®

Current Network Status
 Last Updated: Sat Sep 21 17:14:14 UTC 2024
 Updated every 90 seconds
 Nagios® Core™ 4.4.6 - www.nagios.org
 Logged in as nagiosadmin

Host Status Totals

Up	Down	Unreachable	Pending
3	0	0	0

[All Problems](#) [All Types](#)

Service Status Totals

Ok	Warning	Unknown	Critical	Pending
3	0	0	0	0

[All Problems](#) [All Types](#)

Host Status Details For All Host Groups

Limit Results: 100

Host	Status	Last Check	Duration
app-vm1.fra1.internal	UP	09-21-2024 17:12:17	0d 1h 25m 56s
db-vm1.fra1.db	UP	09-21-2024 17:13:57	0d 1h 22m 25s
web-vm1.fra1.web	UP	09-21-2024 17:10:37	0d 1h 24m 12s

Results 1 - 3 of 3 Matching Hosts

Problems
 Services (Unhandled)
 Hosts (Unhandled)
 Network Outages

Quick Search:

1) Prepare a docker-compose for a nginx server.

Requirements:

- nginx logs need to survive between nginx container restarts
- docker should use network bridge subnet 172.20.8.0/24

```
version: '3.8'
```

```
services:
```

```
  nginx:
```

```
    image: nginx:latest
```

```
    volumes:
```

```
      - nginx_logs:/var/log/nginx
```

```
    networks:
```

```
      my_bridge:
```

```
        ipv4_address: 172.20.8.2
```

```
    ports:
```

```
      - "80:80"
```

```
volumes:
```

```
  nginx_logs:
```

```
    driver: local
```

```
networks:
```

```
  my_bridge:
```

```
driver: bridge
ipam:
  config:
    - subnet: 172.20.8.0/24
```

```
root@localhost~# docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED        STATUS        PORTS                               NAMES
42973b6b08e8   nginx:latest  "/docker-entrypoint..."  50 seconds ago  Up 11 seconds  0.0.0.0:80->80/tcp, :::80->80/tcp  root-nginx-1
```

```
[root@localhost ~]# docker network inspect root_my_bridge
```

```
[root@localhost ~]# docker network inspect root_my_bridge
[
  {
    "Name": "root_my_bridge",
    "Id": "fba38e429322f6dba3bcf0d5206e186bf2f6eaf5c9a9b16fdb06b4fb632e0db9",
    "Created": "2024-09-22T00:08:24.292357801+06:00",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "172.20.8.0/24"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {
      "42973b6b08e8e4fa99fd9b2902b86d8e6dc832d5dd674bcd43a42dabcc956d2e": {
        "Name": "root-nginx-1",
        "EndpointID": "625d4c45a1340e6a2c7c64f4ef0ab98c7cdcaff746e9d64a17e7b73f79a45162",
        "MacAddress": "02:42:ac:14:08:02",
        "IPv4Address": "172.20.8.2/24",
        "IPv6Address": ""
      }
    },
    "Options": {},
    "Labels": {
      "com.docker.compose.network": "my_bridge",
      "com.docker.compose.project": "root",
      "com.docker.compose.version": "2.20.2"
    }
  }
]
```

1) Which Kubernetes command you will use to identify the reason for a pod restart in the project "internal" under namespace "production".

```
kubectl describe pod my-pod -n production
```

1) Consider the followings:

POD MEMORY(bytes)	NAME	CPU(cores)	
java-app-7d9d44ccbf-lmvbc	java-app	3m	951Mi
java-app-7d9d44ccbf-lmvbc	java-app-logrotate	1m	45Mi
java-app-7d9d44ccbf-lmvbc	java-app-fluentd	1m	84Mi
java-app-7d9d44ccbf-lmvbc	mongos	4m	62Mi

Application pod has the following resource quota:

- Memory request & limit: 1000 & 1500
- CPU request & limit: 1000 & 2000
- Xmx of 1000M

Java-app keep restarting at random. From Kubernetes configuration perspective, what are the possible reasons for the pod restarts?

Possible Reasons for Pod Restarts:

If the pod exceeds its memory limit, the Linux kernel will invoke the Out-Of-Memory (OOM) killer, causing the pod to restart with the reason OOMKilled.

Please use the accompanied elasticsearch helm template to create a Kubernetes deployment of elasticsearch.

Provide a screenshot & deployment yaml of the resultant deployment in Kubernetes.

```
root@master-node:/home/dnet/elasticsearch_helm/elasticsearch/templates# kubectl describe pod elasticsearch-0 -n elasticsearch
Name:          elasticsearch-0
Namespace:     elasticsearch
Priority:       0
Service Account: default
Node:          worker01/192.168.3.40
Start Time:    Sun, 22 Sep 2024 11:49:02 +0000
Labels:        app=elasticsearch
               controller-revision-hash=elasticsearch-7b9477c4d
               statefulset.kubernetes.io/pod-name=elasticsearch-0
Annotations:   <none>
Status:        Running
IP:            10.244.1.85
IPs:           IP: 10.244.1.85
Controlled By: StatefulSet/elasticsearch
Init Containers:
  configure-sysctl:
    Container ID:  containerd://506eb85689b460af1a5f9166e673ca890d5209d831038885b42853a21fb198ca
    Image:         docker.elastic.co/elasticsearch/elasticsearch:7.17.0
    Image ID:      docker.elastic.co/elasticsearch/elasticsearch@sha256:577b382dda5d05385aea8c7b60dad97e02ff41ca0da54f723151c2aed9ac8f54
    Port:          <none>
    Host Port:     <none>
    Command:
      sysctl
      -w
      vm.max_map_count=262144
    State:         Terminated
      Reason:       Completed
      Exit Code:    0
      Started:      Sun, 22 Sep 2024 11:49:05 +0000
      Finished:     Sun, 22 Sep 2024 11:49:05 +0000
    Ready:         True
    Restart Count:  0
  Limits:
    cpu:          25m
    memory:       128Mi
  Requests:
    cpu:          25m
    memory:       128Mi
  Environment:    <none>
```

<https://github.com/zeesaan/DevOps.git>

Explain how Prometheus work

Prometheus collects data in the form of time series through a pull model. The Prometheus server scrapes a list of data sources (often called exporters) at a specified polling frequency.

Data is stored as metrics, each identified by a unique name used for referencing and querying. Prometheus stores data locally on disk, which allows for fast data storage and querying, but it also has the capability to store metrics in remote storage systems. Each Prometheus server is standalone, operating independently without relying on network storage or other remote services.

How do you create custom Prometheus alerts and alerting rules for Kubernetes monitoring? Provide an example alert rule and its configuration.

```
groups:
  - name: kubernetes-alerts
    rules:
      - alert: HighCpuUsage
        expr:
sum(rate(container_cpu_usage_seconds_total{job="kubelet", cluster="",
image!="", container!="POD"}[5m])) by (pod) > 0.8
        for: 5m
        labels:
          severity: warning
        annotations:
          summary: "High CPU usage detected"
          description: "Pod {{ $labels.pod }} is using more than 80%
CPU."
```

What is the Prometheus query you can use in Grafana to properly show usage trend of an application metric that is a counter?

Grafana using Prometheus, use the following query:

```
rate(your_counter_metric_name[5m])
```

Create a New Dashboard in Grafana.

Add a Panel and select Graph.

Choose Prometheus as the data source.

Enter your query (e.g., `rate(http_requests_total[5m])`).

Save your dashboard.